


Looking for edges graph




Gaussian derivative (speed things up):

- don't have to differentiate whole thing -diff gaussians * f

- x high -> low
- y low -> high

Take laplacian of gaussian (2nd derivative - zero crossing)



- more sigma smooths it out over wider -> less edges in more detailed areas
- 1. smoothing
- 2. edge enhancement
- 1&2 can be done through derivative of gaussian
- Canny edge detection**
- 3. edge localization (threshold) -> binary image
 - 1. pixel less than t set to 0
 - 2. others to 1
- non max suppression edge width-> single pixel
 - specific location by gradient direction
- low and high threshold (high starts low finishes edge curves)

Derivatives

- smoothing *negative signs* used to get high response in regions of high contrast
- sum to 0 -> no response in constant regions
- high abs at points of high contrast edge strength = gradient magnitude
- choose min seam

$M(i, j) = Energy(i, j) + \min(M(i-1, j-1), M(i-1, j), M(i-1, j+1))$

$$p \equiv \begin{bmatrix} f & 0 & u_0 \\ 0 & f & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} R_{3 \times 3} & t_{3 \times 1} \\ X_{4 \times 1} \end{bmatrix}$$
$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \equiv \begin{bmatrix} X \\ Y \\ Z \end{bmatrix}$$

Pairs of $[X, Y, Z]$ and $[u, v]$ → eqns to constrain M
How do I get $[X, Y, Z]$, $[u, v]$?

$$p_i \times M X_i = 0$$
$$error = (a^*u + b^*v + c) / \sqrt{a^2 + b^2}$$
$$M = \begin{bmatrix} m_1^T \\ m_2^T \\ m_3^T \end{bmatrix}$$
$$\begin{bmatrix} u_1 \\ v_1 \\ 1 \end{bmatrix} \times \begin{bmatrix} m_1^T X_1 \\ m_2^T X_1 \\ m_3^T X_1 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

Harris Corner Detection

- Need:
- **repeatability:** despite geometric transformations - correspondence
 - Saliency: small distinctive subset
 - compactness/opt: fewer features than pixels
 - locality: small area of image (no clutter)
 - How:
 - eigenvalues give us basis vector
 - eigenvalue >> other -> edge
 - rotation invariant: by eigen to find new basis rotation
 - scale invariant: no -> yes
 - choosing window size: expand radius -> get most activation for each pixel

1. cornerness: M matrix
2. threshold: keep high
3. local filter for most cornerlike
 - 1. non-max suppression
 - Formulas:
 - $R = \det(M) - \alpha \text{trace}(M)^2$
 - R small: flat, $R > 0$: corner,

Blob detection

- edge = ripple, (2nd of gauss) -> blob is where ripples dip lower
 - mag of laplacian will provide scale of blob - smaller laplacian dip wants smaller blob -> (characteristic scale)

Stereo & Camera Calibration

Given 2 cams + correspondence → find coord for real world point
Estimate intrinsic matrix K
 $[u \ v \ 1]^T = [K \ t] [x \ y \ z]^T$
Use $u, v \leftrightarrow x, y, z$ correspondences to calculate

Linear Method estimate K:
 $p_i = \lambda M X_i, \lambda \neq 0 \rightarrow p_i \times M X_i = 0$
Cross product → skew symmetric

$$\begin{bmatrix} 0^T & -X_i^T & v_i X_i^T \\ X_i^T & 0^T & -u_i X_i^T \\ -v_i X_i^T & u_i X_i^T & 0^T \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

11 degs. freedom, ambig. to scale
 $\argmin \|A n\|_2 \text{ s.t. } \|n\|_2 = 1 \rightarrow$
 $e_vec \text{ of } A^T A \text{ w/ small } e_val$
In practice:
1. Get $K[R \ t]$ from $M \rightarrow$
QR decomp of leftmost 3x3 in $M \rightarrow$ upper triangular matrix * rotational
2. Extra correspondences = over constrained
So → set up linear equations
Non-linear opt (ie least squares):
$$\sum \|proj(M X_i) - [u_i, v_i]^T\|_2^2$$

Triangulation: solve for X given M:
Geom: midpoint of viewing rays

- trying to do multiple scales
- simple descriptors: raw pixel vectorized (highly sensitive to noise/shifting) -> SIFT
 - Histograms to bin pixels sub-patches according to their gradient orientation (don't vectorize)
 - mag: old determines weight of histo
 - angle: tan
 - compute histograms for different patches

Orient patch by max weight and make this the 0

- less bins -> less orientations = less info
- more bins -> too much info

- partially invariant to
- illumination changes
 - camera viewpoint
 - clutter

- $M = X \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} X^T$
- find patches that have most similar (lowest SSD)
 - $(X_i - X_j)^2 + (Y_i - Y_j)^2$ for every distinct pair (i, j)
 - robustness: distance to best match/distance to second best match
 - =1 is ambiguous
 - lowest: first match looks good

• hypothesize transformation -> apply trans and see if more match

Convolution

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i-u, j-v]$$

$$G = H * F$$

Cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i+u, j+v]$$

$$G = H \otimes F$$
$$\min \|Lh\|^2 \text{ s.t. } \|h\|^2 = 1$$

• h are unknowns (h1 is row)

• equations: $x' = \frac{h_1^T p}{h_1^T p}$

• minimize: $x' - \frac{h_1^T p}{h_1^T p} = 0$

$$L = \begin{bmatrix} p_1^T & 0 & -x'_1 p_1^T \\ 0 & p_2^T & -y'_1 p_2^T \\ \vdots & \vdots & \vdots \end{bmatrix}$$

• $L =$


• $\min(h^T A h) \text{ s.t. } \|h\|^2 = 1$

• smallest eigenvalue solves

- 4x4 matrix of histograms of pixel gradient angles
Rotate histogram to dominant orientation
Invariant to: rotation, scale ~camera&env. change:
Find matches: find two SIFT with lowest sum of squares distance (SSD), take closest (1 - k)
Robustness r.t.: dist to best / dist to 2nd best
→ use matches to hypothesize transformation T

Clustering algorithms

1. want to find centers
 - minimize SSD between all points in cluster
2. methods
 - k means: randomly initialize k clusters -> for each point find closest c -> given points avg for c -> if c changed iterate again (without random)
 - gives local min (converges reasonably in time), sensitive to start/initial clusters/must cluster all points(outliers suck), detects spherical clusters
 - Feature space
 - intensity+position or color
 - normalized cuts: have similar appearance forming parts of an object
 - build graph node for every pixel -> every edge has weight which is similarity (affinity)



- $Ncut(A, B) = \frac{cut(A, B)}{assoc(A, V)} + \frac{cut(A, B)}{assoc(B, V)}$
- $\exp(-1/(2\sigma^2)) * ||x_i - x_j||^2$
 - delete links that cross btw segments
 - cut low affinities
 - min cut (summing up all the weights taht you cut) - removal makes graph disconnected
 - generalized eigenvalue problem
 - pro: does not require model fo data distribution, flexible choice of affinities
 - cons: comp high, dense, preference fore balanced partitions (equal weights - or else could favor really small clusteters ex: small object on big background)

Reducing noise = average across img.
Correlation Filtering: (G=new, F=old, H=filter)

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i+u, j+v]$$

Smoothies by evening out spike

in matrix. (Flips filter bot to top, right to left)

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix} = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

Blur with Box Sharpening Filter Shift Left by 1 pi.
-< Sharpens the image

Properties of Convolution →

Grassman's Law

1. If two test lights can be matches with the same set of weights, then they match each other
2. If we scale the test light, then the matches get scaled by the same amount.
3. If we mix two test light, then mixing the matches will match the result (superposition)

Linear Color Space

RGB = Single Wavelength primaries, good for devices but not for perception
CIEXYZ = Y value approximates brightness, project to display: (x, y) = (x/x+y+z), (y/x+y+z)
Non-Linear Color Space
HSV = Hue, Saturation (Purity, intensity), Value-Nonlinear-reflects topology of colors by coding hue as an angle. More User friendly than the other two.
Distance in color space is a problem that no one is really trying to solve except for **LAB**

Sampling

- sample the 2d space on grid -> quantize each sample (int)
 - one value per pixel
- sample across R, G, B (can be avged together for one image)
Filter (denoise, resize, extract texture, edges, detect patterns)
- enhance image - denoise
 - raw pixel of same image won't be same
- Salt and pepper: white pixels - gaussian noise sample from gaussian N(u, sigma) more var = more noise
Reducing noise: average of neighbors - expect neighbors to be similar $[1, 1, 1, 1, 1]/5$ (uniform) $[1, 4, 5, 4, 1]/16$

Correlation filtering: $G[i, j] = 1/(2k+1)^2$
 $\sum_{u=-k}^k \sum_{v=-k}^k F[i+u, j+v]$ G = HxF H is mask produces flipped -> convolution then apply cross (symmetric will output same)

- full = any part of g touches f, same = same size as f, valid = doesn't fall of edge
- boundary (clip filter black, wrap around, copy edge, reflect across edge)
- window size doesn't impact gaussian -but variance does directly (filter size ~ 6σ) choose window size from this
- remove hf; low pass filter

Runtime: $O(n^2 m^2)$ -> 1d G * 1d G = 2d G
 $O(N)$ (by seperability outer product)
prop: convolution is linear, can shift
Non-linear filters: median (comp heavy)
low freq: smoothing
high freq: og + (og - smooth = details)

$$Filter_{ij} \propto \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$
$$Filter_{ij} \propto \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{x^2}{2\sigma^2}\right) \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{y^2}{2\sigma^2}\right)$$

Start with two points (x_i, y_i)
 $y = Av$
 $\begin{bmatrix} y_1 \\ y_2 \end{bmatrix} = \begin{bmatrix} x_1 & 1 \\ x_2 & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix}$
 x_1, y_1
 $\|y - Av\|^2 = \left\| \begin{bmatrix} y_1 \\ y_2 \end{bmatrix} - \begin{bmatrix} mx_1 + b \\ mx_2 + b \end{bmatrix} \right\|^2 = (y_1 - (mx_1 + b))^2 + (y_2 - (mx_2 + b))^2$
Roberts
 $M_x = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$ $M_y = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$
Prewitt
 $M_x = \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}$ $M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$
Sobel
 $M_x = \begin{bmatrix} 1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$ $M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & 2 & -1 \end{bmatrix}$

$$\frac{\partial f(x, y)}{\partial x} \begin{bmatrix} 1 & -1 \end{bmatrix} \text{ Partial Deriv } \frac{\partial f(x, y)}{\partial y} \begin{bmatrix} -1 & 1 \end{bmatrix}$$

Larger σ = larger scale edges detected
Smaller σ = fine features detected

